

Hadoop

Types of Data

The following three types of data can be identified:



Structured data:

Data which is represented in a tabular format

E.g.: Databases



Semi-structured data:

Data which does not have a formal data model

E.g.: XML files



Unstructured data:

Data which does not have a pre-defined data model

E.g.: Text files

5V'S OF BIG DATA

Velocity

Speed at which data is emanating and changes are occurring between the diverse data sets

Velocity



Volume

This refers to the sheer volume of data being generated every second.

Volume



Variety

Variety

Can use structured as well as unstructured data.



Veracity

Veracity

Data reliability and trust.
Verifying and validating the data



Value

Value

Having access to big data is all well and good but that's only useful if we can turn it into a value.



Big Data Storage & Computation ?



Storing Big Data was a Problem

*Even if a part of Big Data is Stored-
Processing it, took Years*



Hadoop Solves Big Data Problems



Storing Big Data was no more a Problem



And Processing did not take Years

Write Once



Data will be written to the HDFS once and then read several times

Write Once



Updates to files after they
have already been closed
are not supported

HDFS Architecture


NameNode and DataNode

NameNode

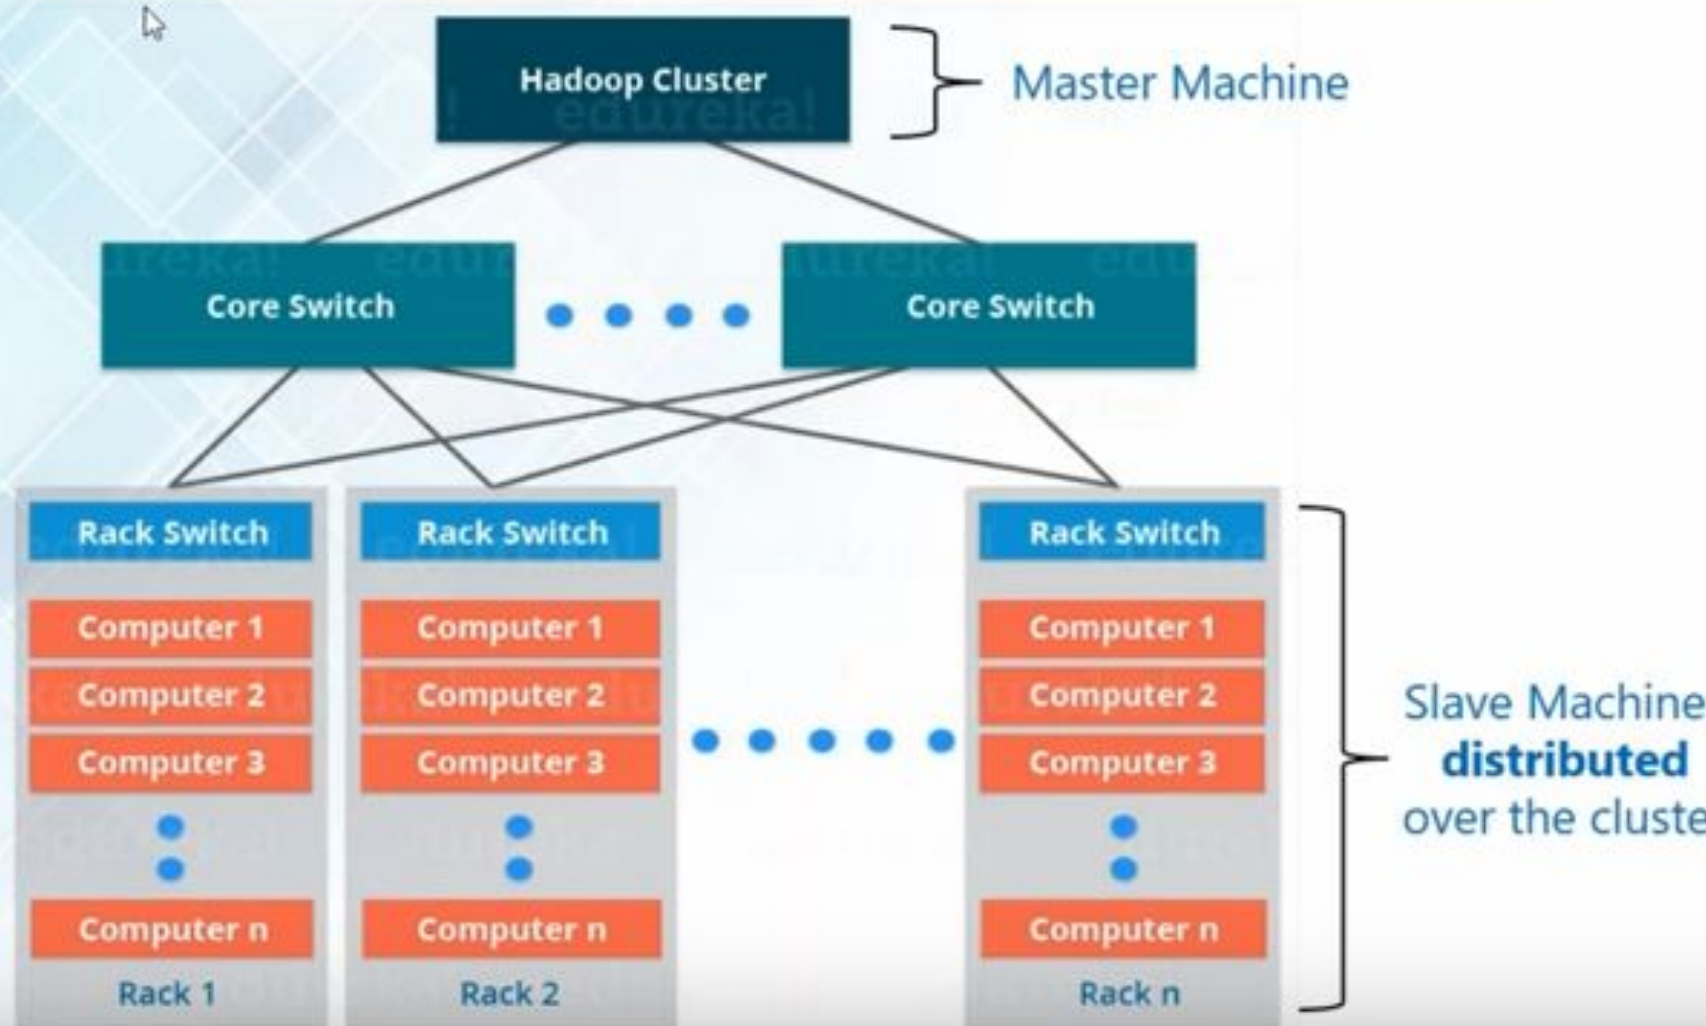
- Master daemon
- Maintains and Manages DataNodes
- Records metadata e.g. location of blocks stored, the size of the files, permissions, hierarchy, etc.
- Receives heartbeat and block report from all the DataNodes

DataNode

- Slave daemons
- Stores actual data
- Serves read and write requests from the clients



Hadoop Cluster Architecture – Master Slave Topology



Hadoop cluster




Cluster of machine running Hadoop at Yahoo!

HDFS block

- Each file is stored on HDFS as blocks
- The default size of each block is **128 MB** in Apache Hadoop 2.x (64 MB in Apache Hadoop 1.x)
- Let us say, I have a file example.txt of size 248 MB. Below is the representation of how it will be stored on HDFS




How many blocks will be created if a file of size 514 MB is copied to HDFS ?



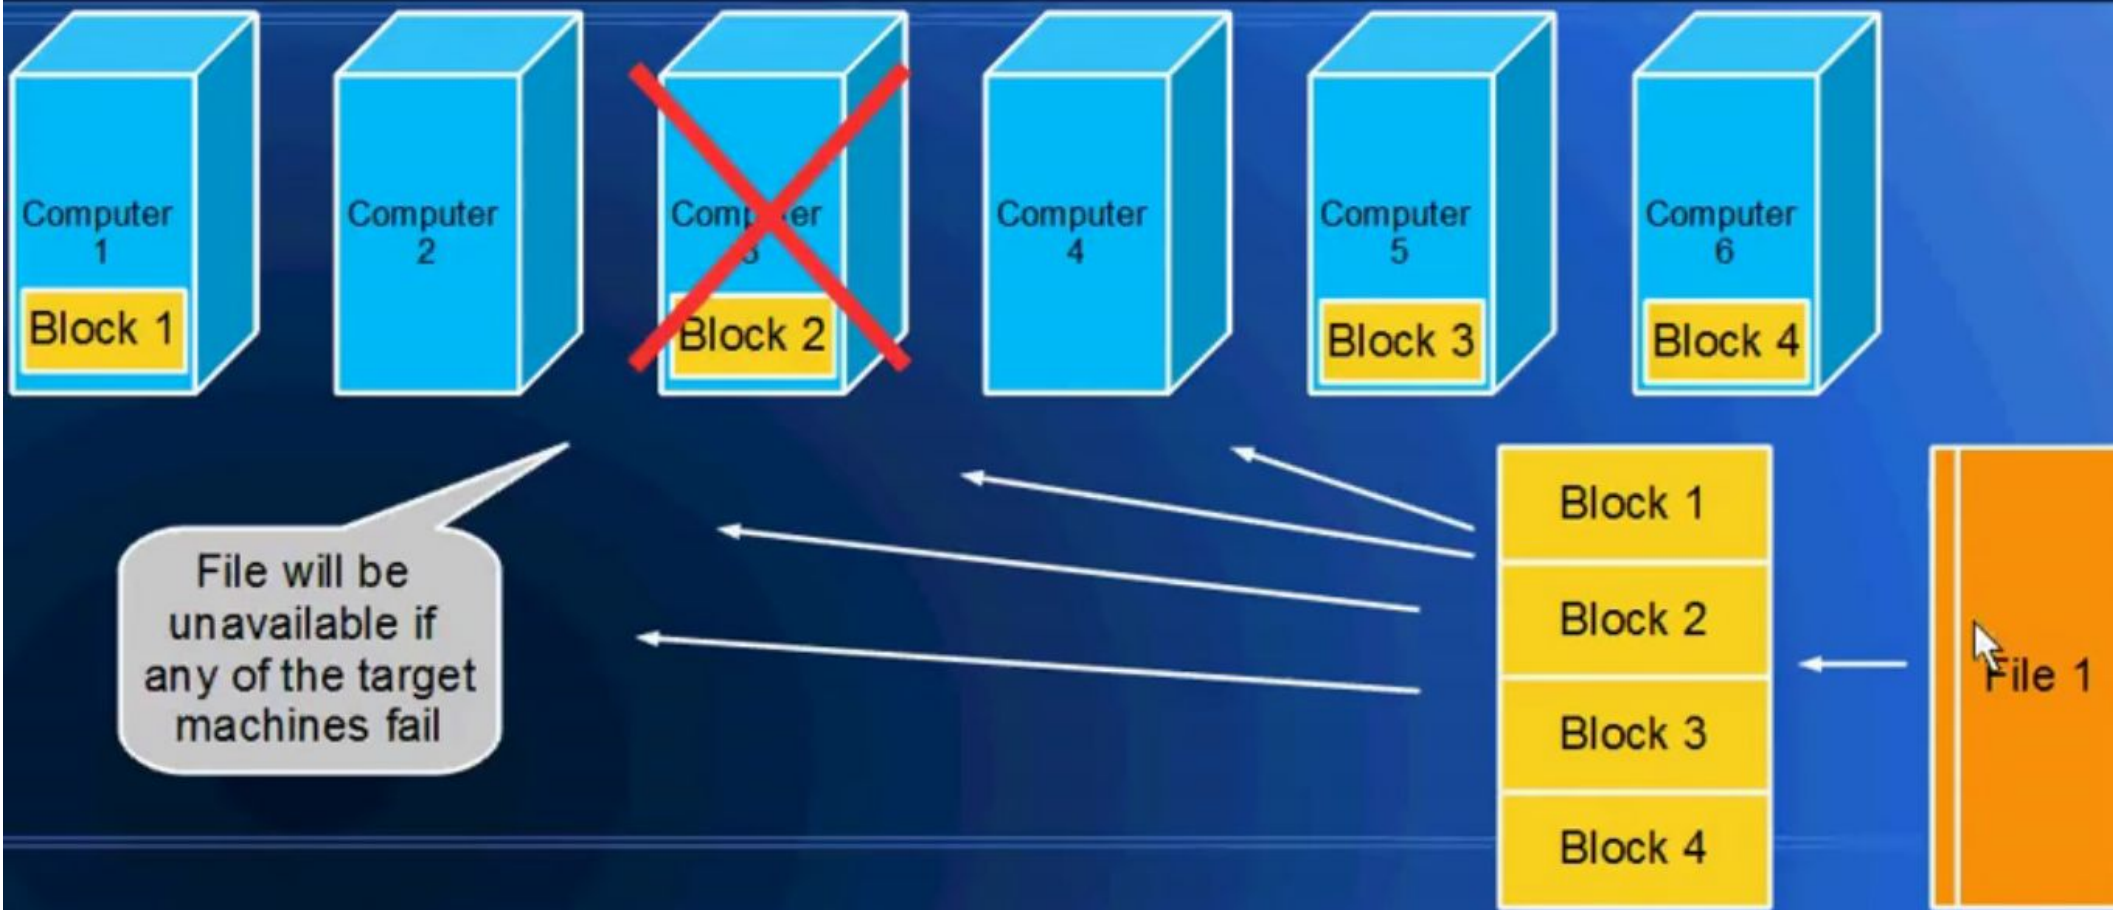
Target Machines

Cluster




Hardware Failure

Cluster







Performance Lost in Failures



Master

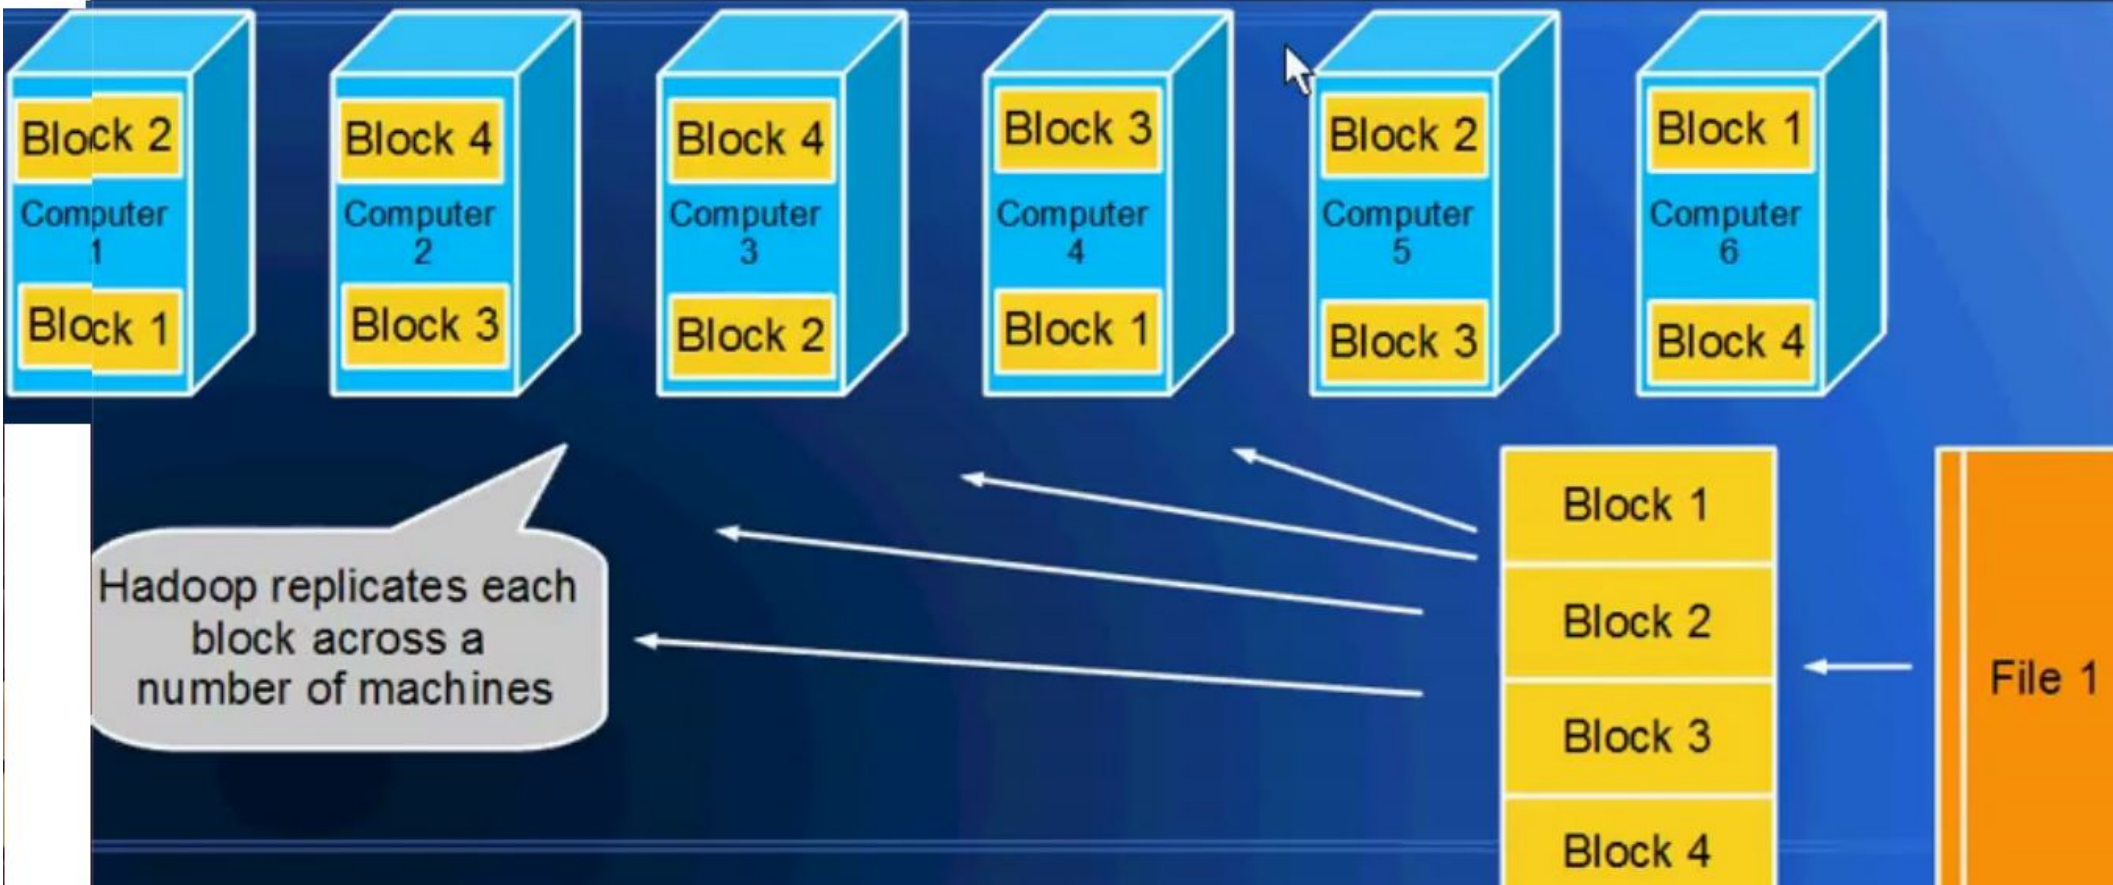
Slaves

System performance lost
is in proportion to the
number of nodes failed




Replication

Cluster

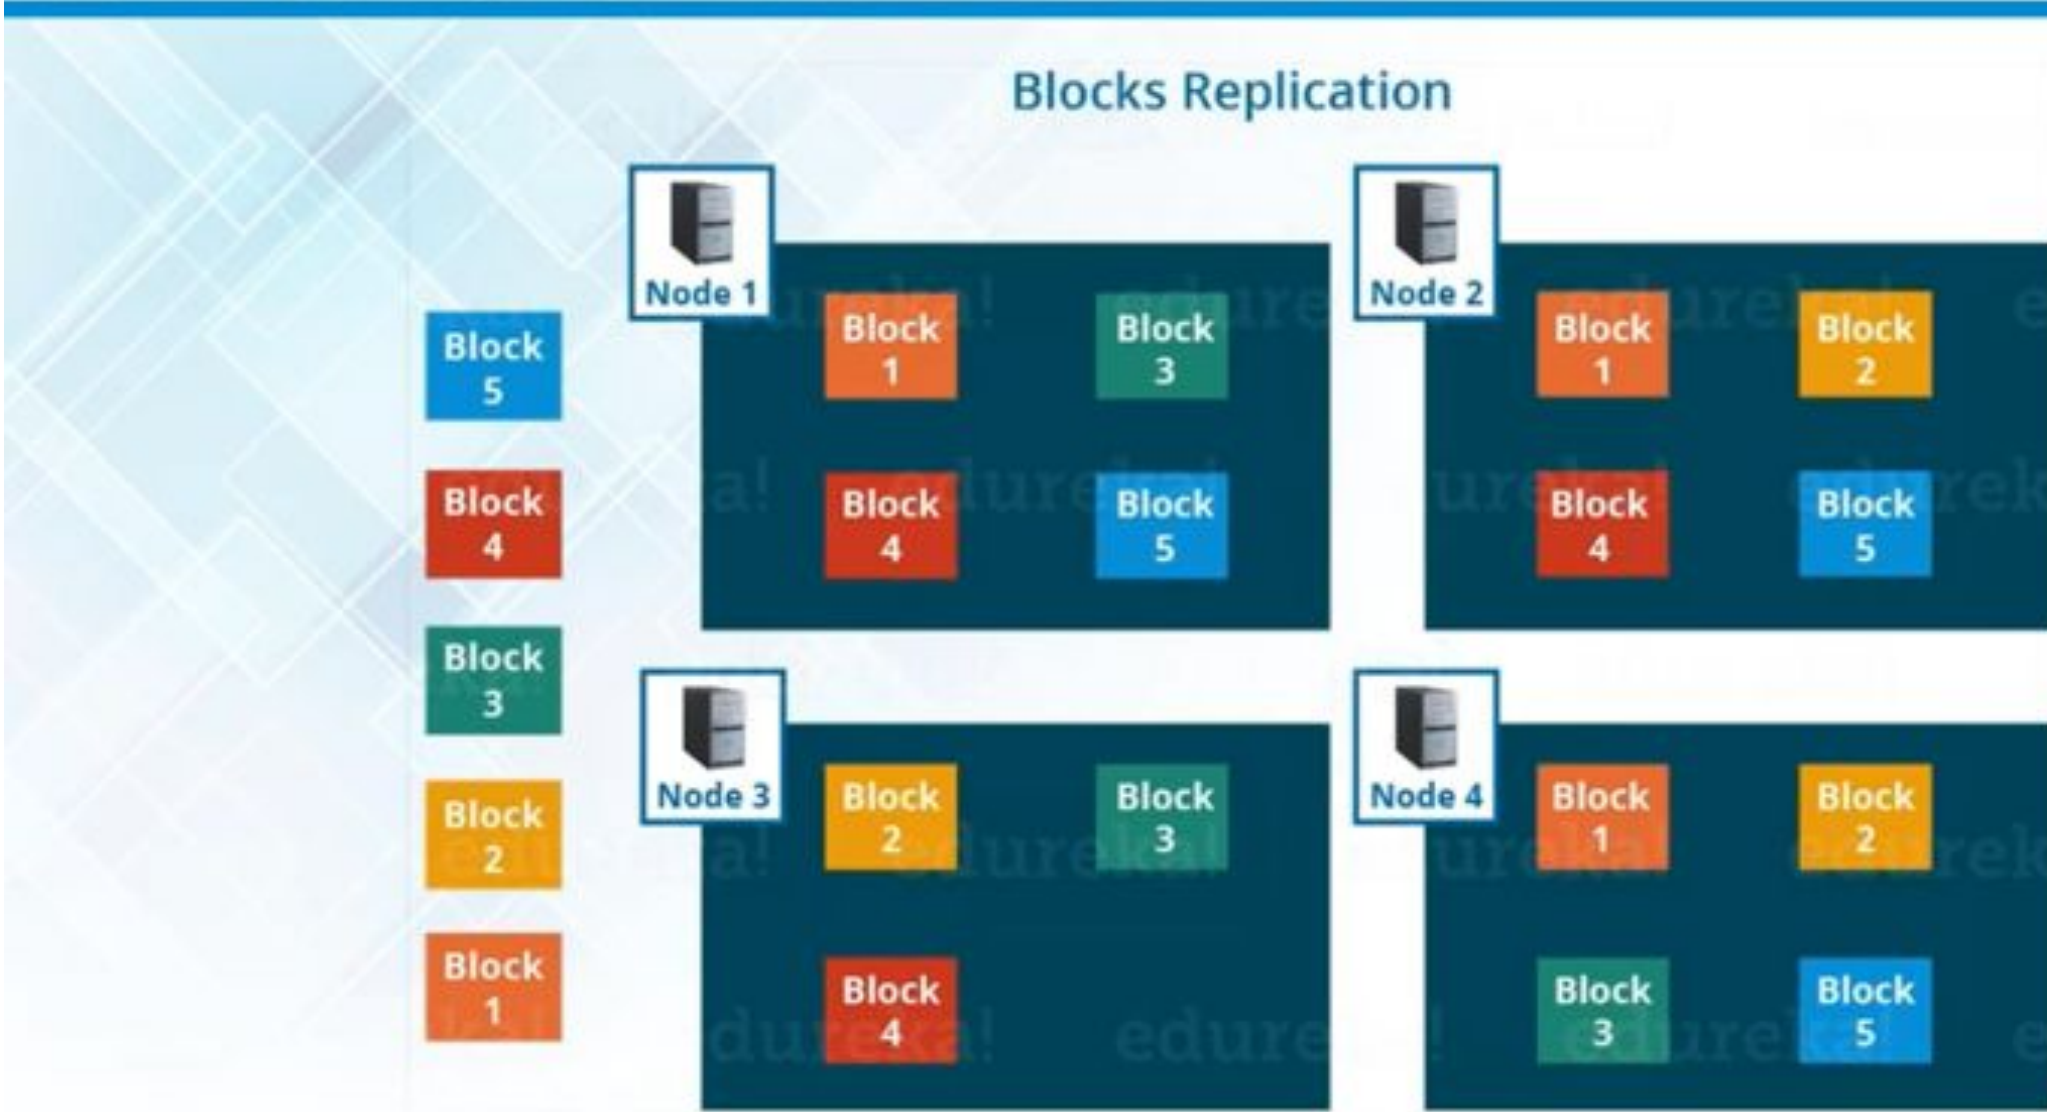


Number of Copies

Cluster



Hadoop Architecture – Block Replication



Hadoop Architecture: Rack Awareness

Rack Awareness Algorithm

Block A : 

Block B: 

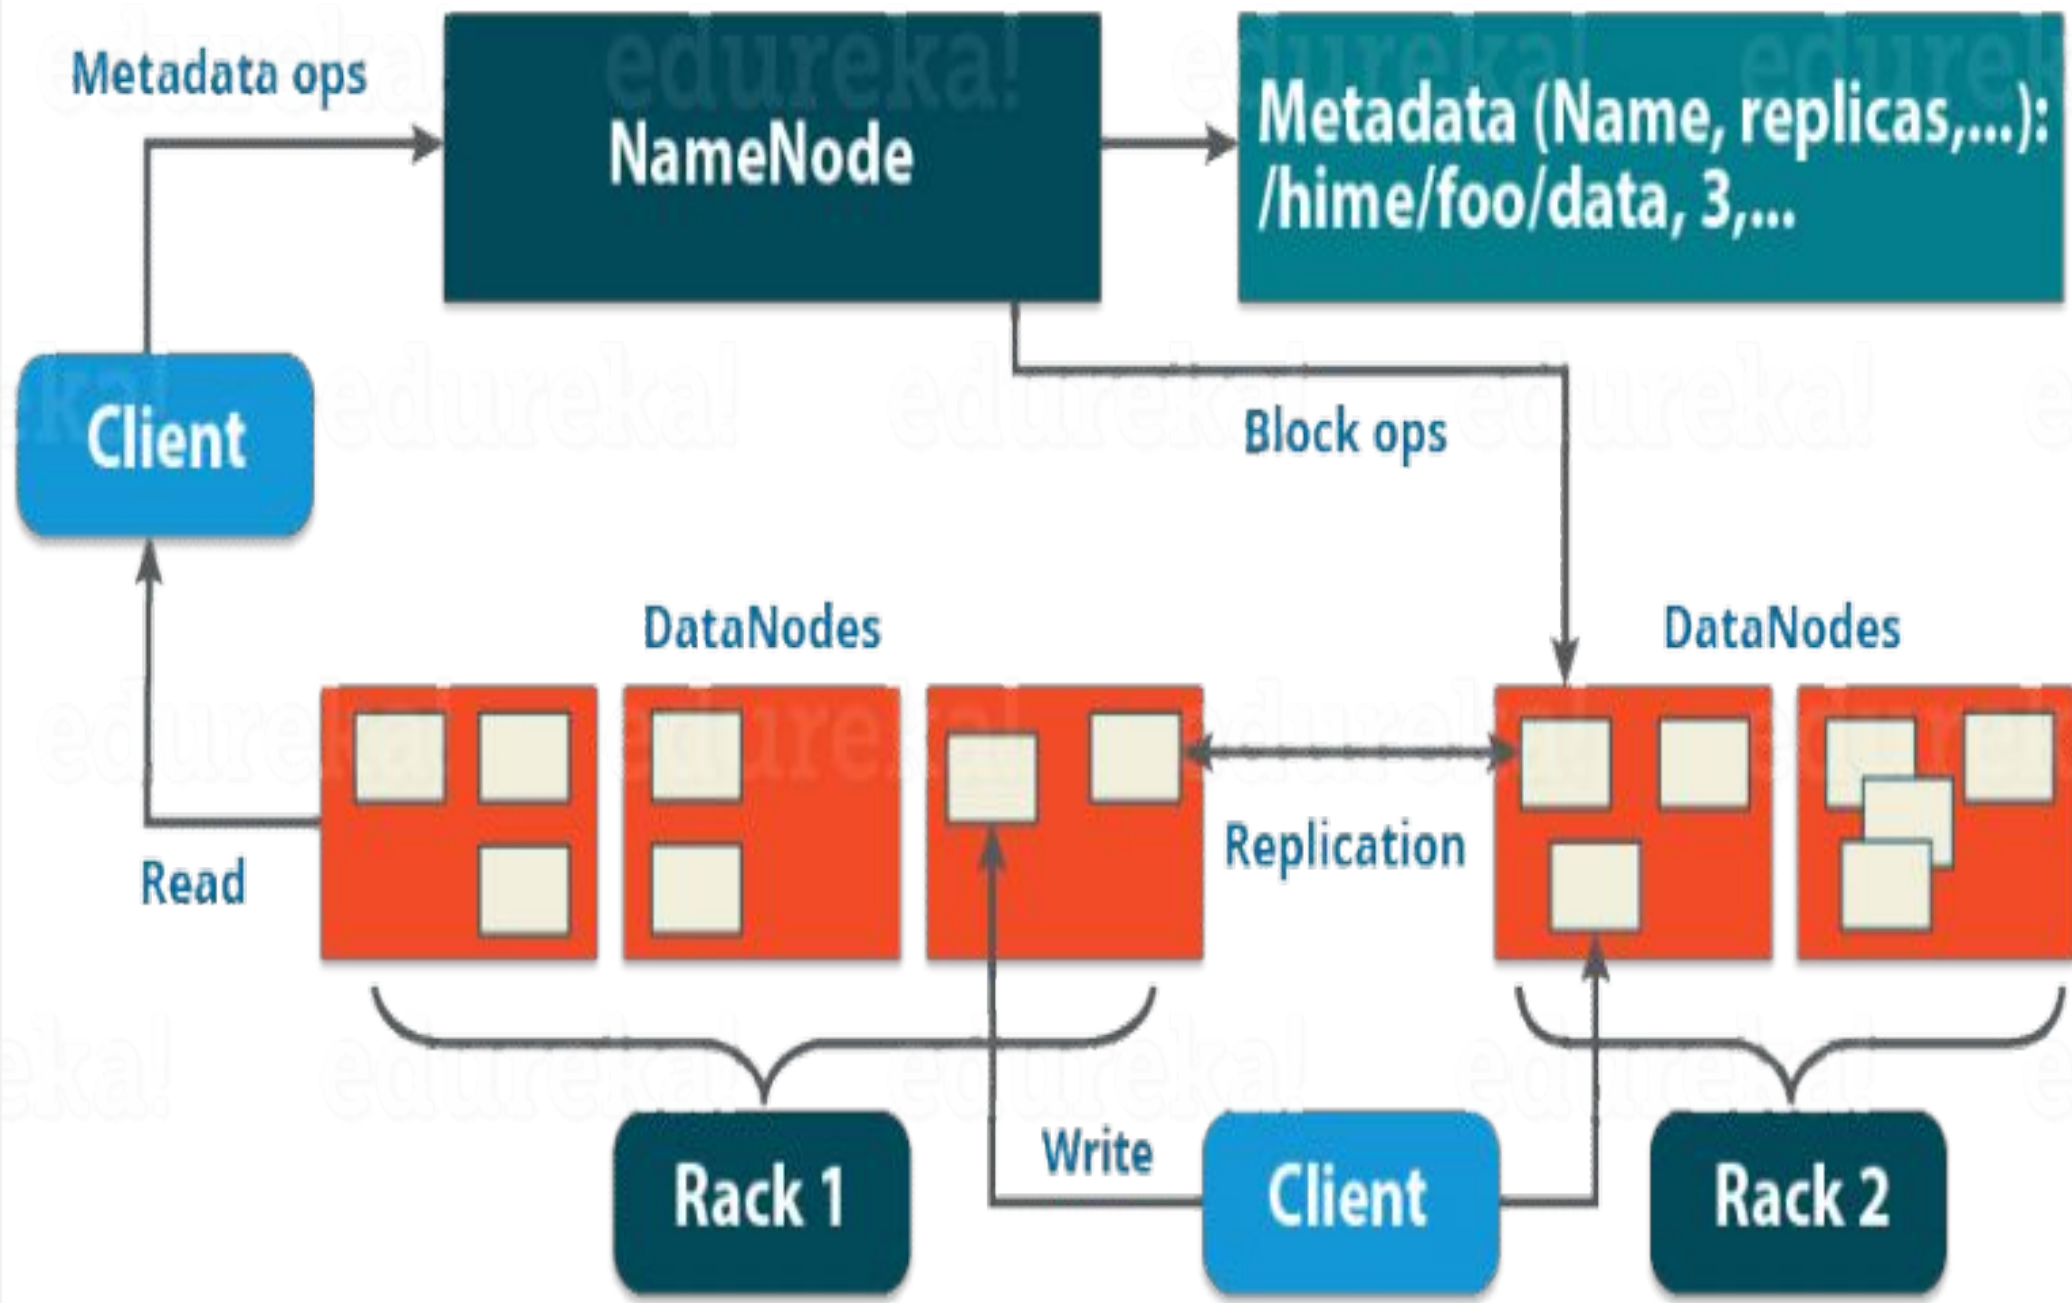
Block C: 



Suggested replication topology

- 1st replica placed on same node as client;
- 2nd replica placed on different rack from 1st rack; and
- 3rd replica placed on same rack as 2nd rack, but on a different node.

HDFS Architecture



HDFS Operation Principle

The HDFS components comprise different servers like NameNode, DataNode, and Secondary NameNode.

NameNode Server (single instance)

- Maintains the file system name space
- Manages the files and directories in the file system tree
- Stores information in the namespace image and the edit log
- NameNode knows the data nodes on which all the blocks for a given file exist
- NameNode is a critical one point failure node

DataNode Server (multiple instances)

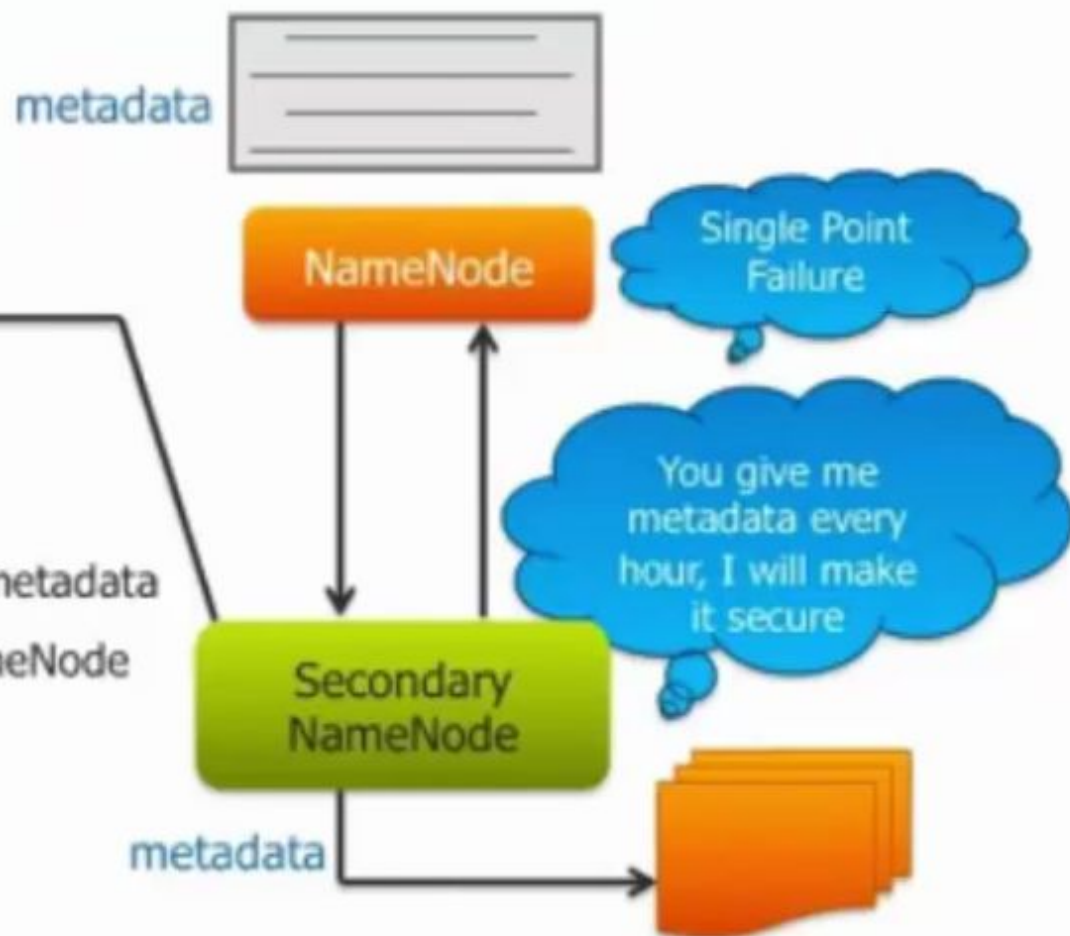
- Associated with data storage places in the file system
- Reports to NameNode periodically with lists of blocks they store
- Stores and retrieves blocks when referred by clients or NameNode
- Servers read, write requests, performs block creation, deletion, and replication upon instruction from NameNode

Secondary NameNode Server (single instance)

- Not exactly a hot backup of the actual NameNode server
- Used for recovery of NameNode in case of NameNode failure
- Keeps namespace image through edit log periodically
- Namespace image lags behind, so total recovery is impossible

Secondary Name Node

- ✓ **Secondary NameNode:**
 - ✓ Not a hot standby for the NameNode
 - ✓ Connects to NameNode every hour*
 - ✓ Housekeeping, backup of NameNode metadata
 - ✓ Saved metadata can build a failed NameNode




Hadoop processing Unit

Map Reduce


How MapReduce Works

Hadoop Processing Architecture



Map Reduce
Layer

HDFS
Layer



- The client submits the job to Job Tracker
- Job Tracker asks Name node the location of data
- As per the reply from name node, the Job Tracker ask respective task trackers to execute the task on their data
- All the results are stored on some Data Node and the Name Node is informed about the same
- The task Trackers inform the job completion and progress to Job Tracker
- The Job Tracker inform the completion to client
- Client contacts the Name Node and retrieve the results

Counting Word Frequencies

Consider a large text file

Twinkle twinkle little star
How I wonder what you are
Up above the world so high
Like a diamond in the sky
Twinkle twinkle little star
How I wonder what you are
.....



How?

Word	Frequency
above	14
are	20
how	21
star	22
twinkle	32
...	..

Twinkle twinkle little star
How I wonder what you are
Up above the world so high
Like a diamond in the sky
Twinkle twinkle little star
How I wonder what you are
.....

MapReduce Flow

**The raw data is really large
(potentially in PetaBytes)**

**It's distributed across many
machines in a cluster**

**Each machine holds a partition of
data**

MapReduce Flow

Twinkle twinkle little star
How I wonder what you are



Up above the world so high
Like a diamond in the sky



Twinkle twinkle little star
How I wonder what you are




Each partition is given to a different process i.e. to mappers



MapReduce Flow

Twinkle twinkle little star


How I wonder what you are



**Each mapper
works in parallel**


Up above the world so high

Like a diamond in the sky




Twinkle twinkle little star

How I wonder what you are



Map Flow

Twinkle twinkle little star
How I wonder what you are



Within each mapper, the rows are processed serially

Map Flow

Twinkle twinkle little star
How I wonder what you are

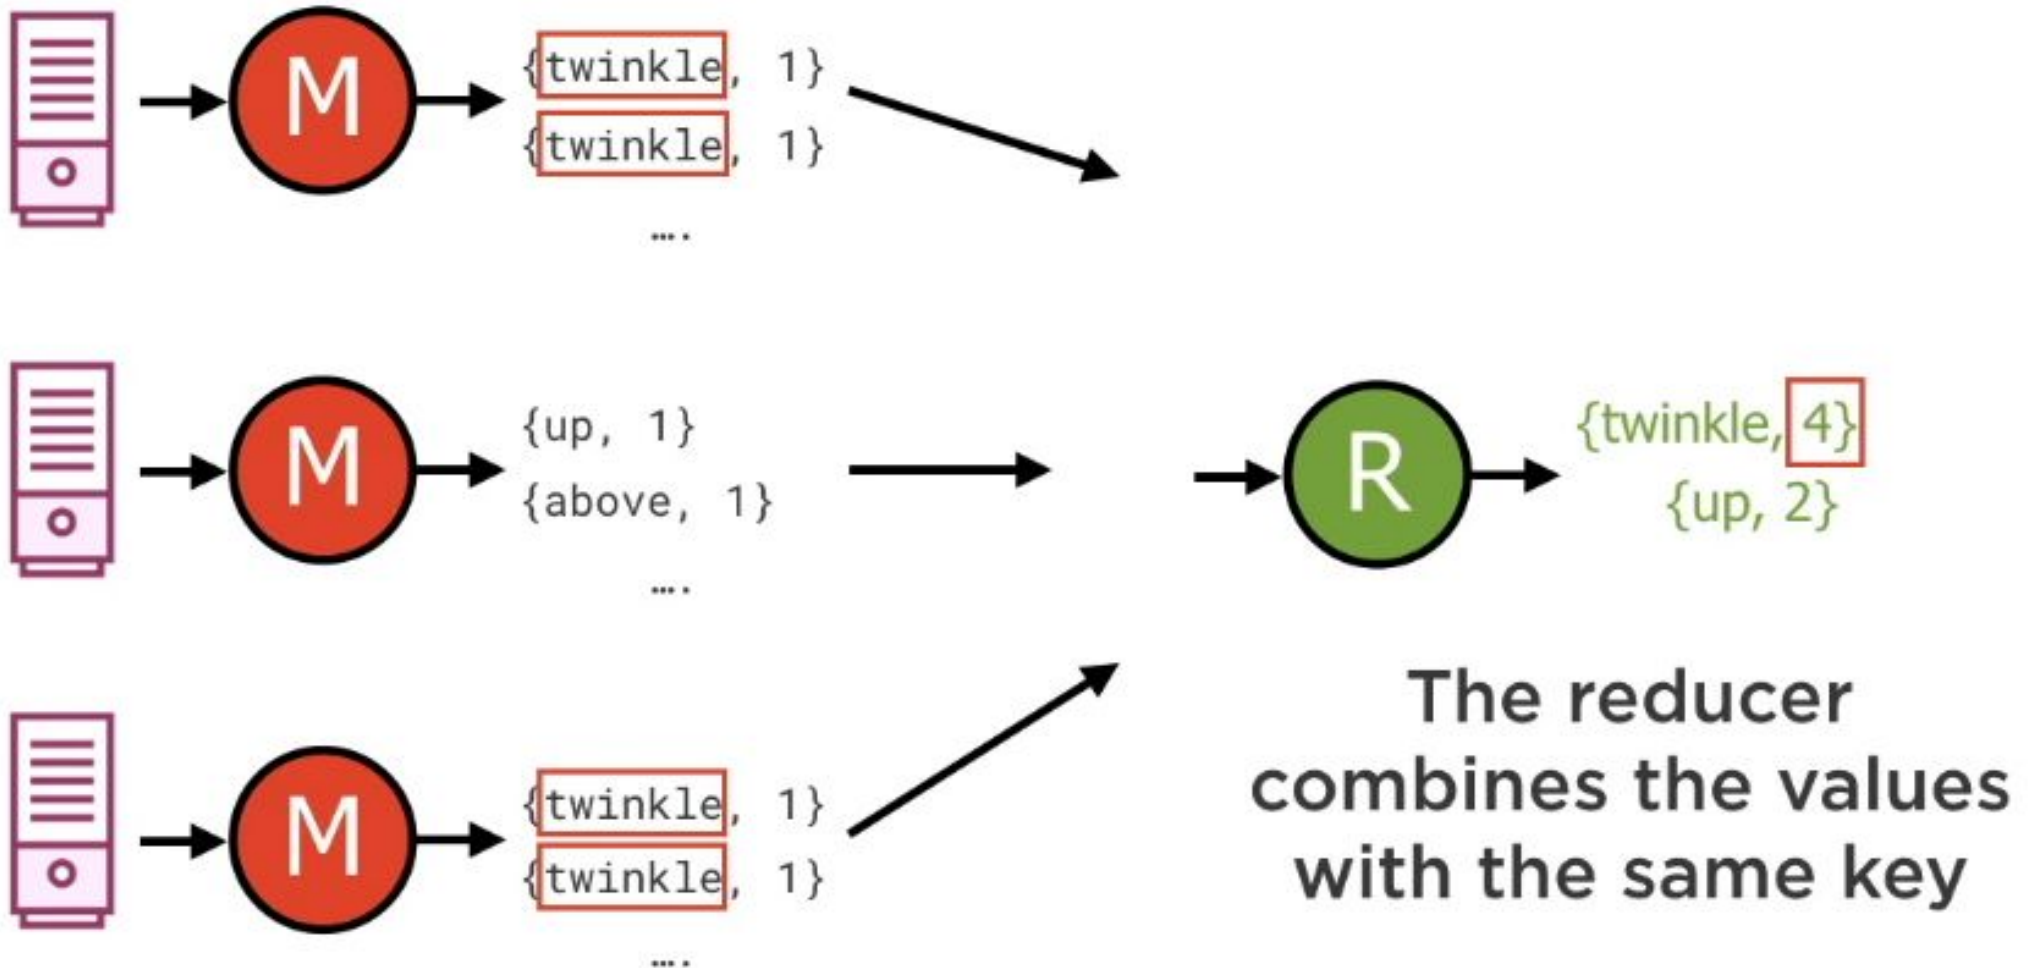


Word	# Count
------	---------




{twinkle, 1}
{twinkle, 1}
{little, 1}
{star, 1}

Each row emits {key, value} pairs

Reduce Flow



Reduce Flow



Key Insight Behind MapReduce



Many data processing tasks can be expressed in this form

Counting Word Frequencies

Twinkle twinkle little star
How I wonder what you are
Up above the world so high
Like a diamond in the sky



**For each word
in each line**

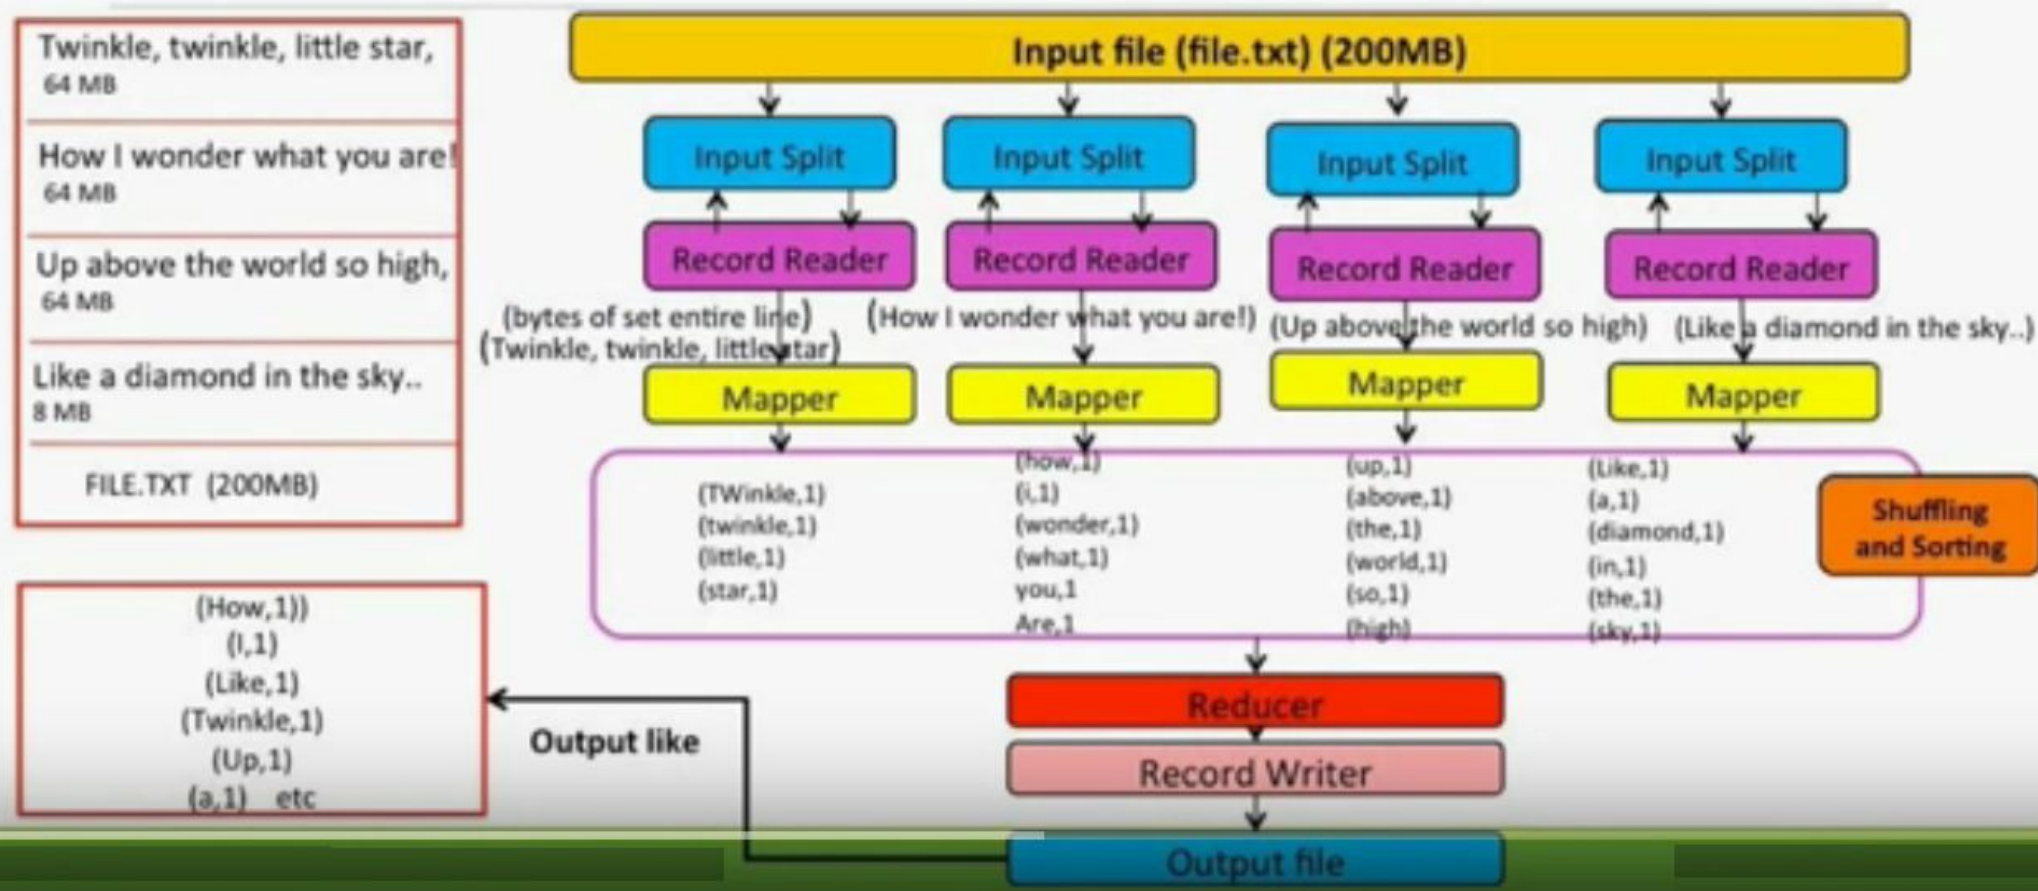
```
{twinkle, 1}  
{twinkle, 1}  
{little, 1}  
{star, 1}
```

..

Word	Count
twinkle	2
little	1
...	...
...	...
...	...
...	...



Internal Process of WordCount Job



Hadoop Framework Keys & Values

➤ JAVA

- int
- long
- float
- double
- boolean
- null
- String

HADOOP

IntWritable

LongWritable

FloatWritable

DoubleWritable

BooleanWritable

NullWritable

Text